

Applied Computer Vision (CS-696)

Homework Assignment 4

Name – Dhaval Harish Sharma

Red ID – 824654344

Question:

The goal of this assignment is to create a local feature matching algorithm using techniques described in Szeliski chapter 4.1. The pipeline we suggest is a simplified version of the famous SIFT pipeline. The matching pipeline is intended to work for instance-level matching -- multiple views of the same physical scene.

Solution:

The three major steps of the local feature matching algorithm are listed and explained as follows:

(1) Interest Point Detection:

In order to create features to match images against, interest points, the position of the features, must be selected. Corners are often good interest points to select since they are different than their neighboring points. To select interest points I used the Harris corner detector algorithm as mentioned in the class notes. The Harris Corner Detection algorithm is explained as follows:

Harris Corner Detection:

The Harris Corner detector works by finding the x and y derivatives of the images. This is accomplished by filtering each image with the Sobel and its transpose. Next the derivatives are squared to form two squared derivatives images and an x-derivative times y-derivative image. The underlying mathematics relies on eigenvalues to describe how corner-like a given patch of an image is, however, this math can be simplified to use the squared derivatives to speed up computations. Pixels that are in a patch with a high corner score are considered distinct and get passed onto the next stage. In order to only operate on single pixels, non-maxima suppression is also run which finds the maximum of patches with scores above the threshold.

(2) Local Feature Description:

After interest points are selected, features must be created from those points. For local feature description, I used SIFT-like features. The steps are as follows:

- (i) Estimate the gradients of the image using filters, one filter for each of the eight gradient directions.
- (ii) For each interest point, create a 4x4 grid around it. Each of the cells contains a histogram made up of the gradients of the pixels in the cell.
- (iii) For simplicity, each pixel only contributes its gradient to the orientation it is closest to.
- (iv) This grid of histograms is flattened into a vector to be used as the feature for an interest point.

(3) Feature Matching:

We match points using the Nearest Neighbor Distance Ratio (NNDR). For each point in image 1, we calculate the proximity of the point's feature descriptor to the descriptors of all the points in image 2. We examine the ratio between the distance to the point "closest" to point A (in terms of the distance of feature descriptors) and the second closest point. If the ratio is sufficiently low (a threshold that we can manipulate as a parameter) then we can safely assume the match is a valid one. The lower the ratio, the higher the

confidence is of the match being a good one. I ended up setting the threshold at about 0.7, as this was not prohibitively low to eliminate all points but was also rather discerning about which matches to accept. A higher threshold will mean more matches will be made, but less accuracy can be guaranteed.

(i) Input: 'features1' and 'features2' are the $n \times$ feature dimensionality features from the two images.

(ii) For each feature in image 1, the nearest neighbor distance ratio is computed with respect to all features in image 2.

(iii) If the ratio is less than a threshold value (0.7), the nearest neighbor of the feature is considered a match and included in the result.

(iv) Sort the matches so that the most confident ones are at the top of the list.

The code files are as follows:

[get_interest_points.m]

```
% Initializing the parameters
alpha = 0.04;
gaussian = fspecial('Gaussian', [25 25], 1);
[gx, gy] = imgradientxy(gaussian);

% Applying the filter to the image
ix = imfilter(image, gx);
iy = imfilter(image, gy);

% Suppress gradients near the edges
ix([(1:feature_width) end-feature_width + (1:feature_width)], :) = 0;
ix(:, [(1:feature_width) end-feature_width + (1:feature_width)]) = 0;
iy([(1:feature_width) end-feature_width + (1:feature_width)], :) = 0;
iy(:, [(1:feature_width) end-feature_width + (1:feature_width)]) = 0;

large_gaussian = fspecial('Gaussian', [25 25], 2);
ixx = imfilter(ix.*ix, large_gaussian);
ixy = imfilter(ix.*iy, large_gaussian);
iyy = imfilter(iy.*iy, large_gaussian);
har = ixx.*iyy - ixy.*ixy - alpha.*(ixx+iyy).*(ixx+iyy);
thresholded = har > 10*mean2(har); % Adaptive threshold

sliding = 1;

switch sliding
    case 0
        components = bwconncomp(thresholded);
        width = components.ImageSize(1);
        x = zeros(components.NumObjects, 1);
        y = zeros(components.NumObjects, 1);
        confidence = zeros(components.NumObjects, 1);
        for ii = 1:(components.NumObjects)
            pixel_ids = components.PixelIdxList{ii};
            pixel_values = har(pixel_ids);
            [max_value, max_id] = max(pixel_values);
            x(ii) = floor(pixel_ids(max_id) / width);
            y(ii) = mod(pixel_ids(max_id), width);
            confidence(ii) = max_value;
        end
end
```

```

case 1
    har = har.*thresholded;
    har_max = colfilt(har, [feature_width feature_width], 'sliding', @max);
    har = har.*(har == har_max);
    [y, x] = find(har > 0);
    confidence = har(har > 0);
end

```

[get_features.m]

```

% Initializing the features
num_points = size(x, 1);
features = zeros(num_points, 128);

% Calculating the small and large gaussian
small_gaussian = fspecial('Gaussian', [feature_width feature_width], 1);
large_gaussian = fspecial('Gaussian', [feature_width feature_width], feature_width / 2);

% Applying the gradient to the image
[gx, gy] = imgradientxy(small_gaussian);
ix = imfilter(image, gx);
iy = imfilter(image, gy);

get_octant = @(x,y) (ceil(atan2(y, x) / (pi / 4)) + 4);

orients = arrayfun(get_octant, ix, iy);
mag = hypot(ix, iy);
c_size = feature_width / 4;

for ii = 1:num_points
    frame_x_range = (x(ii) - 2 * c_size):(x(ii) + 2 * c_size - 1);
    frame_y_range = (y(ii) - 2 * c_size):(y(ii) + 2 * c_size - 1);
    frame_mag = mag(frame_y_range, frame_x_range);
    frame_mag = frame_mag.*large_gaussian;
    frame_orients = orients(frame_y_range, frame_x_range);

    % Looping through each cell in the frame
    for xx = 0:3
        for yy = 0:3
            cell_orients = frame_orients(xx * 4 + 1:xx * 4 + 4, yy * 4 + 1:yy * 4 + 4);
            cell_mag = frame_mag(xx * 4 + 1:xx * 4 + 4, yy * 4 + 1:yy * 4 + 4);
            for o = 1:8
                f = cell_orients == o;
                features(ii, (xx * 32 + yy * 8) + o) = sum(sum(cell_mag(f)));
            end
        end
    end
end

features = diag(1./sum(features, 2)) * features; % Normalize feature vectors

end

```

[match_features.m]

```
% Take a threshold value
threshold = 0.7;
dist_matrix = pdist2(features1, features2, 'euclidean');

[sorted_dist_matrix, indices] = sort(dist_matrix, 2);
inverse_confidences = (sorted_dist_matrix(:, 1) ./ sorted_dist_matrix(:, 2));
confidences = 1 ./ inverse_confidences(inverse_confidences < threshold);

% Matching the features
matches = zeros(size(confidences, 1), 2);
matches(:, 1) = find(inverse_confidences < threshold);
matches(:, 2) = indices(inverse_confidences < threshold, 1);
%>

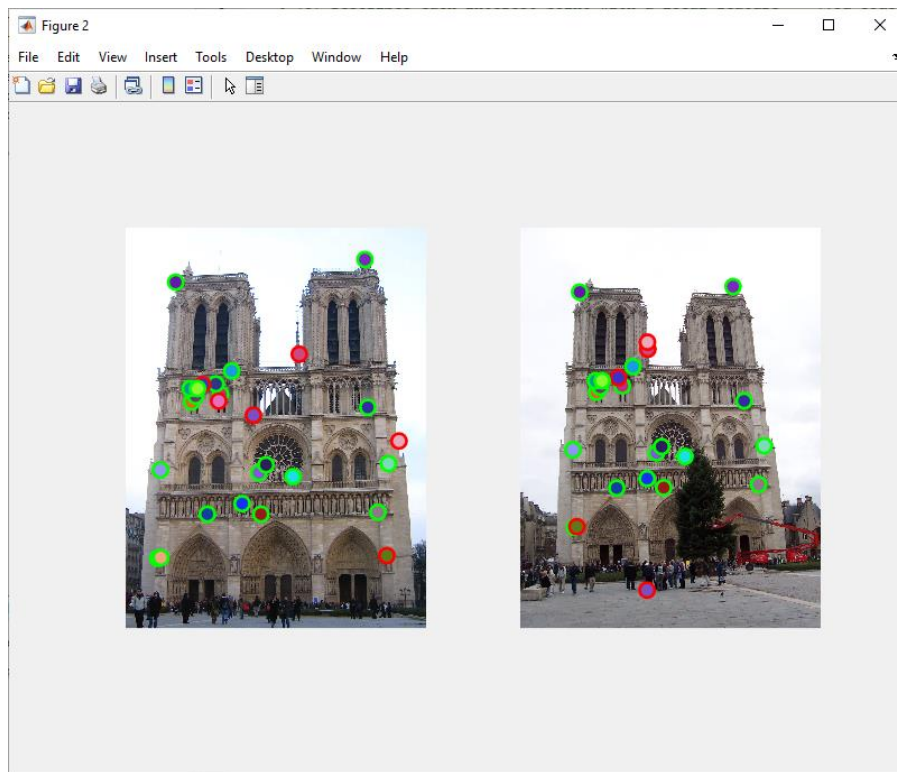
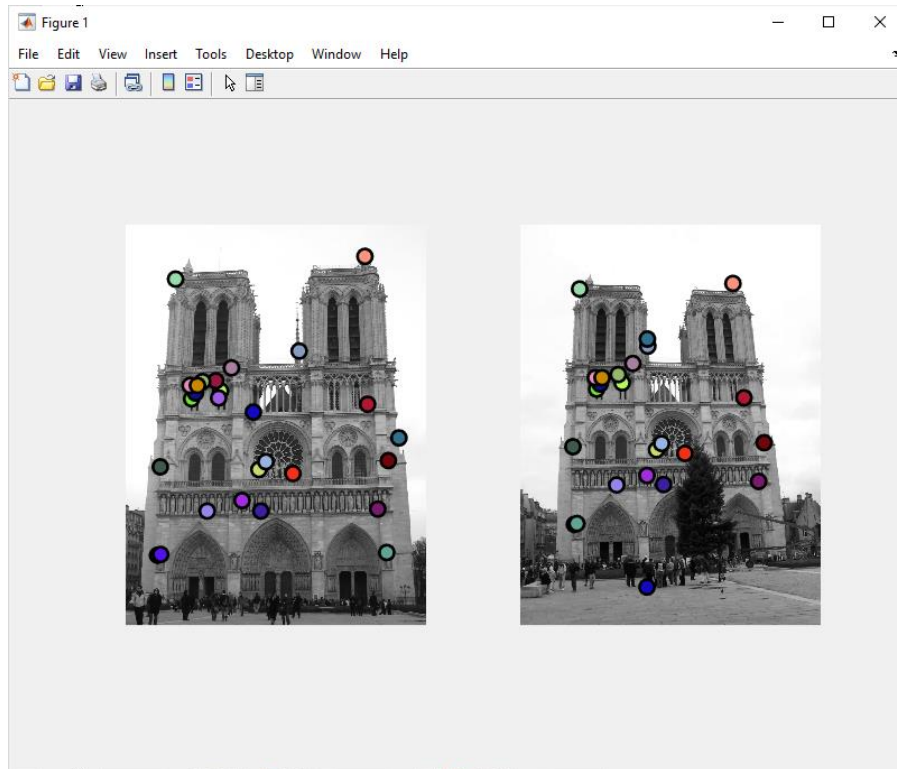
% Sort the matches so that the most confident ones are at the top of the
% list. You should probably not delete this, so that the evaluation
% functions can be run on the top matches easily.
[confidences, ind] = sort(confidences, 'descend');
matches = matches(ind, :);
```

The input images are as follows:



The output of these input images is as follows:

(1) Case 1 (Connected Components):




```
>> proj4
```

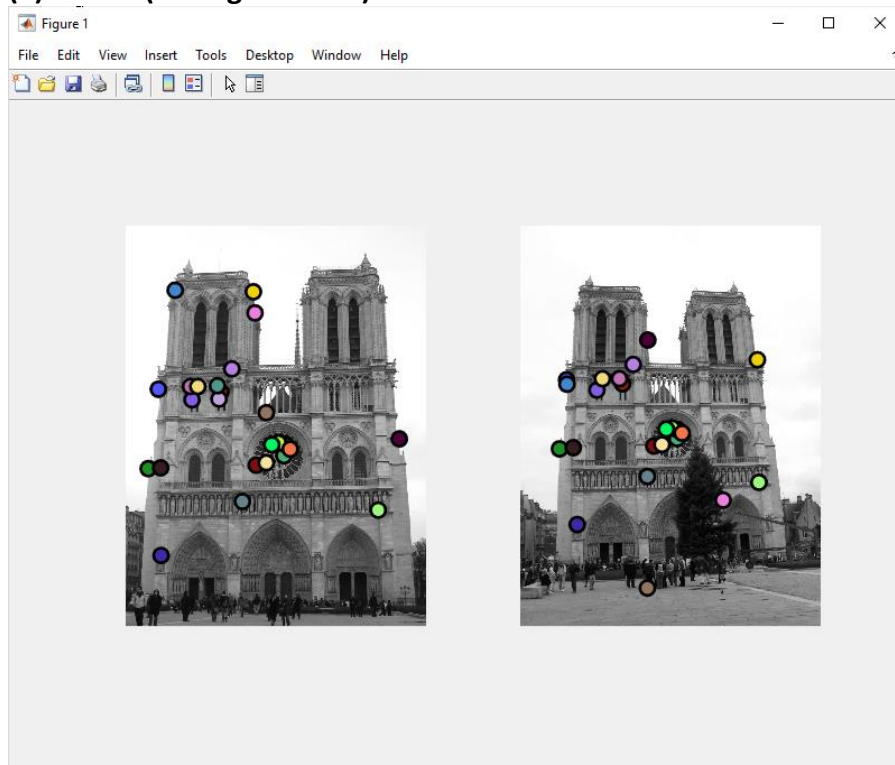
```
Saving visualization to vis.jpg
```

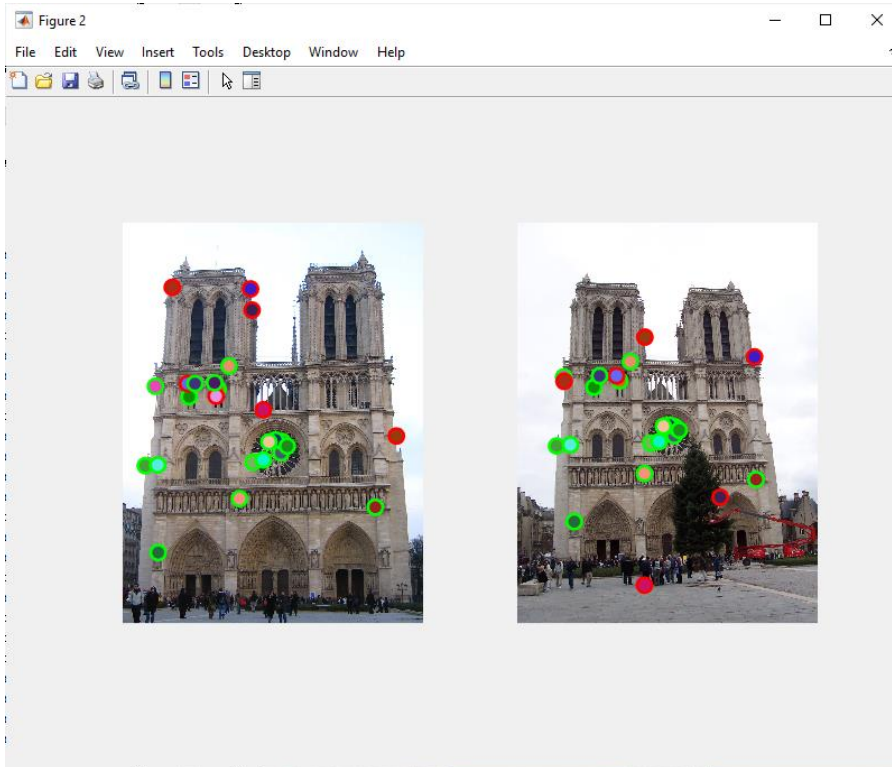
(486, 852) to (518, 804) g.t. point	34 px. Match error	4 px. correct
(164, 1692) to (276, 1522) g.t. point	17 px. Match error	3 px. correct
(476, 888) to (388, 836) g.t. point	4 px. Match error	121 px. incorrect
(338, 892) to (388, 836) g.t. point	3 px. Match error	3 px. correct
(682, 1254) to (692, 1144) g.t. point	67 px. Match error	10 px. correct
(462, 800) to (498, 762) g.t. point	40 px. Match error	1 px. correct
(418, 1466) to (490, 1322) g.t. point	65 px. Match error	5 px. correct
(542, 734) to (572, 706) g.t. point	32 px. Match error	18 px. correct
(654, 960) to (644, 1840) g.t. point	4 px. Match error	942 px. incorrect
(180, 1688) to (288, 1518) g.t. point	5 px. Match error	2 px. correct
(358, 862) to (406, 810) g.t. point	33 px. Match error	7 px. correct
(1344, 1208) to (1238, 1108) g.t. point	56 px. Match error	9 px. correct
(692, 1466) to (728, 1320) g.t. point	11 px. Match error	6 px. correct
(1290, 1456) to (1210, 1304) g.t. point	44 px. Match error	1 px. correct
(330, 824) to (382, 778) g.t. point	66 px. Match error	13 px. correct
(1224, 164) to (1080, 300) g.t. point	19 px. Match error	4 px. correct
(888, 648) to (648, 618) g.t. point	154 px. Match error	196 px. incorrect
(394, 804) to (498, 762) g.t. point	31 px. Match error	68 px. incorrect
(596, 1412) to (644, 1274) g.t. point	56 px. Match error	10 px. correct
(178, 1240) to (268, 1128) g.t. point	16 px. Match error	3 px. correct
(856, 1274) to (836, 1162) g.t. point	19 px. Match error	7 px. correct
(1238, 920) to (1136, 880) g.t. point	20 px. Match error	2 px. correct
(718, 1214) to (718, 1112) g.t. point	24 px. Match error	17 px. correct
(256, 280) to (302, 328) g.t. point	15 px. Match error	4 px. correct
(368, 824) to (416, 778) g.t. point	63 px. Match error	14 px. correct
(1398, 1092) to (646, 582) g.t. point	69 px. Match error	742 px. incorrect
(1336, 1678) to (288, 1518) g.t. point	2 px. Match error	974 px. incorrect

```
21 total good matches, 6 total bad matches
```

```
Saving visualization to eval.jpg
```

(2) Case 2 (Sliding Window):





```
>> proj4
```

```
Saving visualization to vis.jpg
```

(488, 852) to (520, 804) g.t. point	34 px. Match error	4 px.	correct
(810, 1178) to (796, 1082) g.t. point	28 px. Match error	4 px.	correct
(668, 1226) to (678, 1120) g.t. point	46 px. Match error	9 px.	correct
(168, 838) to (236, 784) g.t. point	55 px. Match error	12 px.	correct
(478, 888) to (390, 836) g.t. point	3 px. Match error	121 px.	incorrect
(340, 892) to (390, 836) g.t. point	3 px. Match error	3 px.	correct
(182, 1688) to (290, 1518) g.t. point	5 px. Match error	2 px.	correct
(816, 1120) to (798, 1036) g.t. point	36 px. Match error	8 px.	correct
(1400, 1092) to (648, 582) g.t. point	68 px. Match error	742 px.	incorrect
(782, 1112) to (770, 1028) g.t. point	3 px. Match error	2 px.	correct
(1292, 1456) to (1212, 1304) g.t. point	42 px. Match error	1 px.	correct
(470, 820) to (504, 778) g.t. point	54 px. Match error	4 px.	correct
(116, 1244) to (200, 1134) g.t. point	3 px. Match error	4 px.	correct
(254, 332) to (238, 806) g.t. point	56 px. Match error	455 px.	incorrect
(544, 734) to (574, 706) g.t. point	34 px. Match error	18 px.	correct
(842, 1144) to (822, 1054) g.t. point	55 px. Match error	17 px.	correct
(662, 448) to (1030, 1394) g.t. point	101 px. Match error	974 px.	incorrect
(180, 1240) to (270, 1128) g.t. point	16 px. Match error	3 px.	correct
(720, 958) to (646, 1840) g.t. point	55 px. Match error	941 px.	incorrect
(332, 824) to (504, 778) g.t. point	66 px. Match error	124 px.	incorrect
(654, 340) to (1204, 682) g.t. point	30 px. Match error	645 px.	incorrect
(720, 1214) to (720, 1112) g.t. point	22 px. Match error	17 px.	correct
(748, 1122) to (742, 1034) g.t. point	36 px. Match error	5 px.	correct
(598, 1412) to (646, 1274) g.t. point	55 px. Match error	10 px.	correct
(370, 824) to (418, 778) g.t. point	61 px. Match error	14 px.	correct

```
18 total good matches, 7 total bad matches
```

```
Saving visualization to eval.jpg
```

The other additional input images are as follows:



For these additional images, I have added the ground-truth values. The steps for doing that are as follows:

- (1) Running the `collect_ground_truth_corr()` file. The input of this file will be the above two images.
- (2) After running the file, both the images will be displayed, and we have to mark the corresponded points on

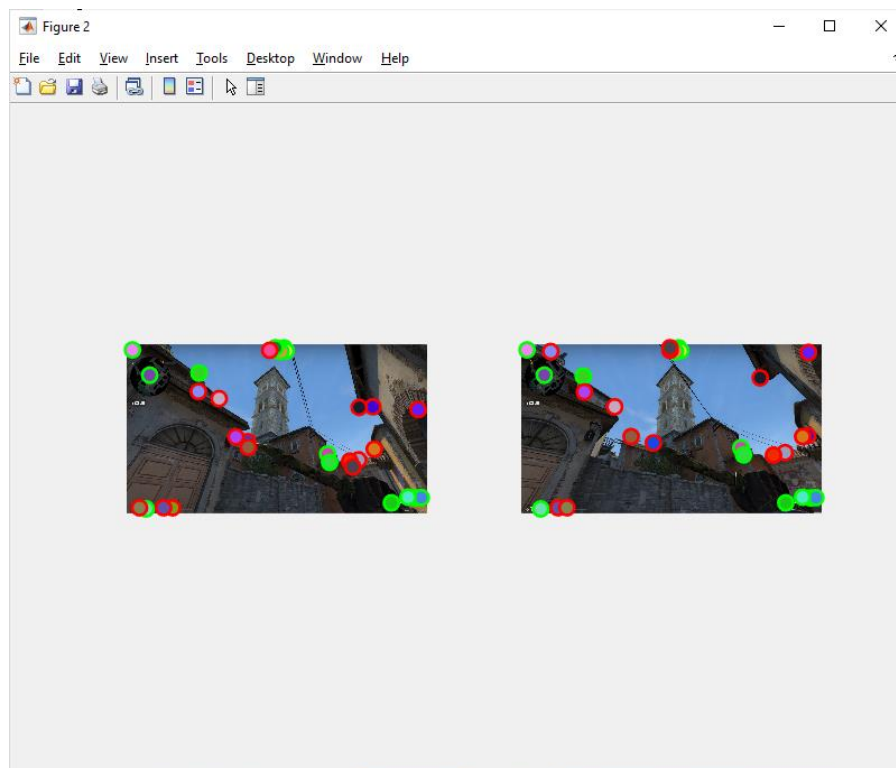
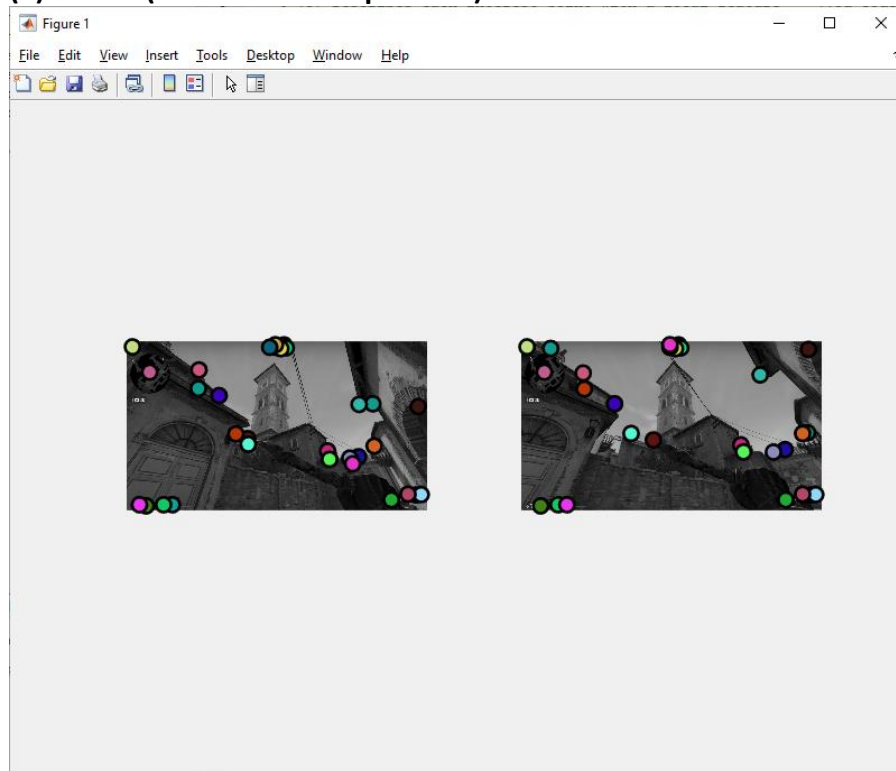
both the images.

(3) After marking the corresponded points we need to click on the negative-ordinate above or left to stop it.

(4) Then the script will be stopped and all the points will be saved in an output file.

(5) We will be using this output file as ground-truth value while running the `evaluate_correspondence.m` file to generate the evaluation image.

(1) Case 1 (Connected Components):



```
>> proj4
```

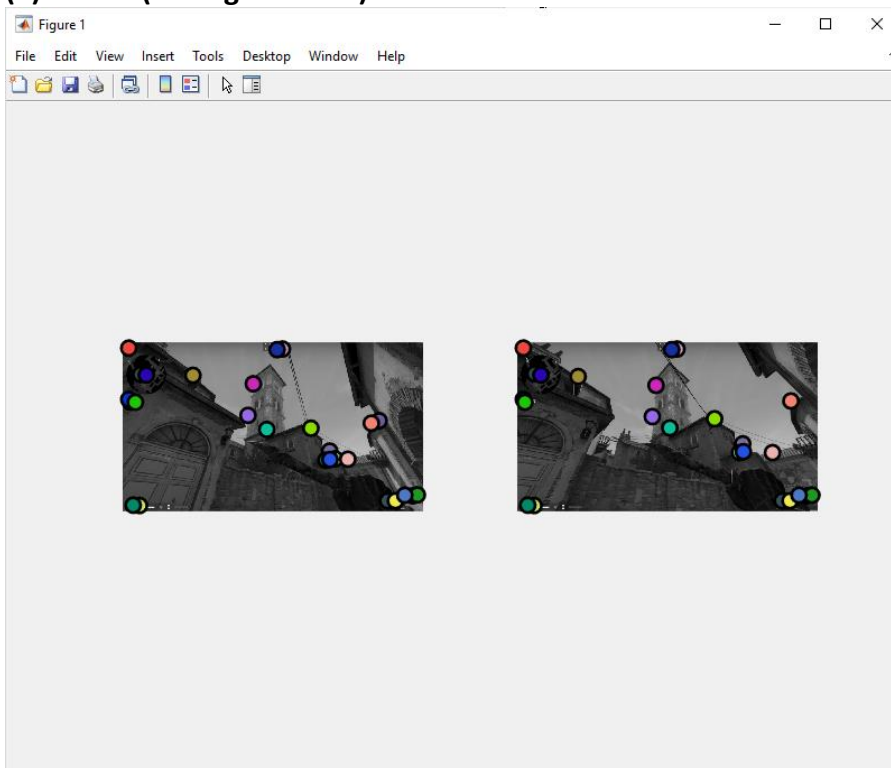
```
Saving visualization to vis.jpg
```

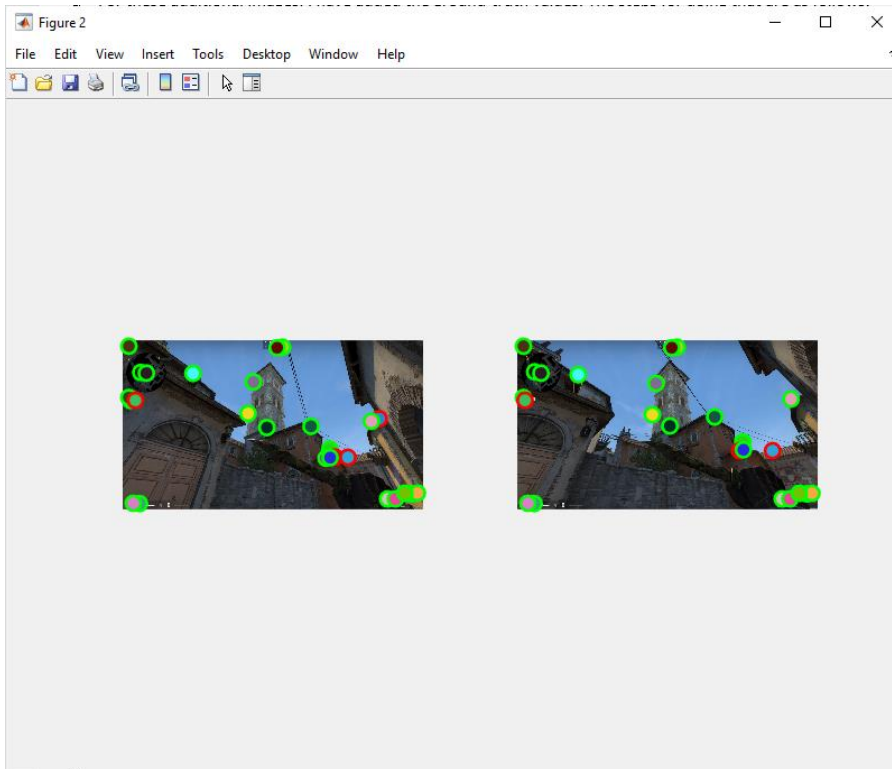
(1020, 44) to (1020, 44)	g.t. point	10 px. Match error	2 px. correct
(1692, 1014) to (1692, 1014)	g.t. point	5 px. Match error	12 px. correct
(986, 48) to (986, 48)	g.t. point	13 px. Match error	2 px. correct
(920, 558) to (974, 548)	g.t. point	4 px. Match error	2 px. correct
(1850, 982) to (1850, 982)	g.t. point	41 px. Match error	6 px. correct
(834, 266) to (886, 276)	g.t. point	24 px. Match error	5 px. correct
(38, 1046) to (38, 1046)	g.t. point	4 px. Match error	3 px. correct
(1486, 404) to (1528, 214)	g.t. point	154 px. Match error	129 px. incorrect
(1638, 502) to (1444, 700)	g.t. point	3 px. Match error	510 px. incorrect
(114, 206) to (110, 202)	g.t. point	8 px. Match error	7 px. correct
(66, 1042) to (66, 1042)	g.t. point	25 px. Match error	3 px. correct
(1882, 976) to (1882, 978)	g.t. point	9 px. Match error	4 px. correct
(698, 600) to (400, 318)	g.t. point	169 px. Match error	465 px. incorrect
(294, 1044) to (292, 1044)	g.t. point	253 px. Match error	5 px. incorrect
(460, 304) to (190, 46)	g.t. point	109 px. Match error	347 px. incorrect
(1322, 696) to (1440, 650)	g.t. point	27 px. Match error	19 px. correct
(1226, 574) to (1320, 530)	g.t. point	45 px. Match error	33 px. incorrect
(800, 468) to (862, 476)	g.t. point	2 px. Match error	3 px. correct
(448, 212) to (388, 220)	g.t. point	17 px. Match error	10 px. correct
(1284, 700) to (598, 406)	g.t. point	55 px. Match error	844 px. incorrect
(1202, 550) to (1260, 490)	g.t. point	12 px. Match error	7 px. correct
(1298, 754) to (1420, 706)	g.t. point	2 px. Match error	3 px. correct
(1802, 978) to (1800, 978)	g.t. point	3 px. Match error	5 px. correct

```
16 total good matches, 7 total bad matches
```

```
Saving visualization to eval.jpg
```

(2) Case 2 (Sliding Window):





```
>> proj4
```

```
Saving visualization to vis.jpg
```

(50, 384) to (50, 384) g.t. point	23 px. Match error	7 px. correct
(40, 366) to (40, 366) g.t. point	3 px. Match error	7 px. correct
(1694, 1014) to (1694, 1014) g.t. point	4 px. Match error	12 px. correct
(1022, 44) to (1022, 44) g.t. point	9 px. Match error	2 px. correct
(988, 48) to (988, 48) g.t. point	15 px. Match error	2 px. correct
(922, 558) to (976, 548) g.t. point	3 px. Match error	2 px. correct
(836, 266) to (888, 276) g.t. point	26 px. Match error	5 px. correct
(108, 1044) to (108, 1044) g.t. point	67 px. Match error	3 px. correct
(116, 206) to (112, 202) g.t. point	7 px. Match error	7 px. correct
(1640, 502) to (1446, 700) g.t. point	1 px. Match error	510 px. incorrect
(40, 38) to (40, 38) g.t. point	9 px. Match error	9 px. correct
(1744, 1014) to (1744, 1014) g.t. point	11 px. Match error	7 px. correct
(1852, 982) to (1852, 982) g.t. point	39 px. Match error	6 px. correct
(1324, 696) to (1442, 650) g.t. point	26 px. Match error	19 px. correct
(1884, 976) to (1884, 978) g.t. point	7 px. Match error	4 px. correct
(1300, 754) to (1422, 706) g.t. point	1 px. Match error	3 px. correct
(1590, 518) to (1750, 376) g.t. point	3 px. Match error	2 px. correct
(1804, 978) to (1802, 978) g.t. point	5 px. Match error	5 px. correct
(152, 210) to (152, 210) g.t. point	35 px. Match error	6 px. correct
(802, 468) to (864, 476) g.t. point	3 px. Match error	3 px. correct
(1358, 746) to (1422, 706) g.t. point	7 px. Match error	68 px. incorrect
(1440, 748) to (1632, 706) g.t. point	44 px. Match error	48 px. incorrect
(68, 1042) to (68, 1042) g.t. point	27 px. Match error	3 px. correct
(450, 212) to (390, 220) g.t. point	18 px. Match error	10 px. correct
(1204, 550) to (1262, 490) g.t. point	13 px. Match error	7 px. correct
(82, 384) to (50, 384) g.t. point	17 px. Match error	26 px. incorrect
(1326, 750) to (1446, 700) g.t. point	7 px. Match error	5 px. correct

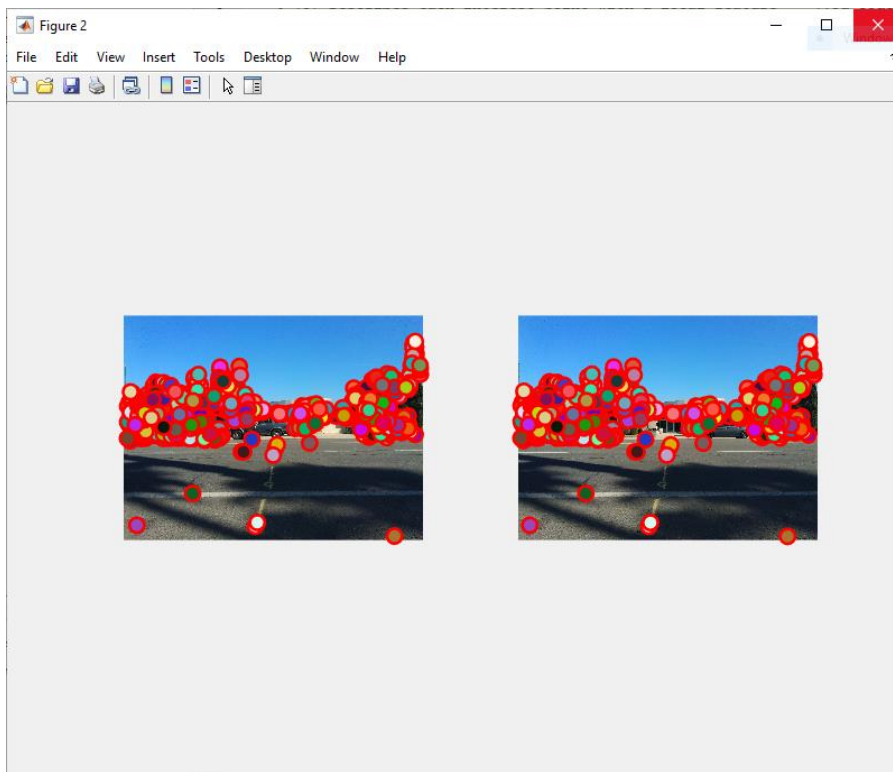
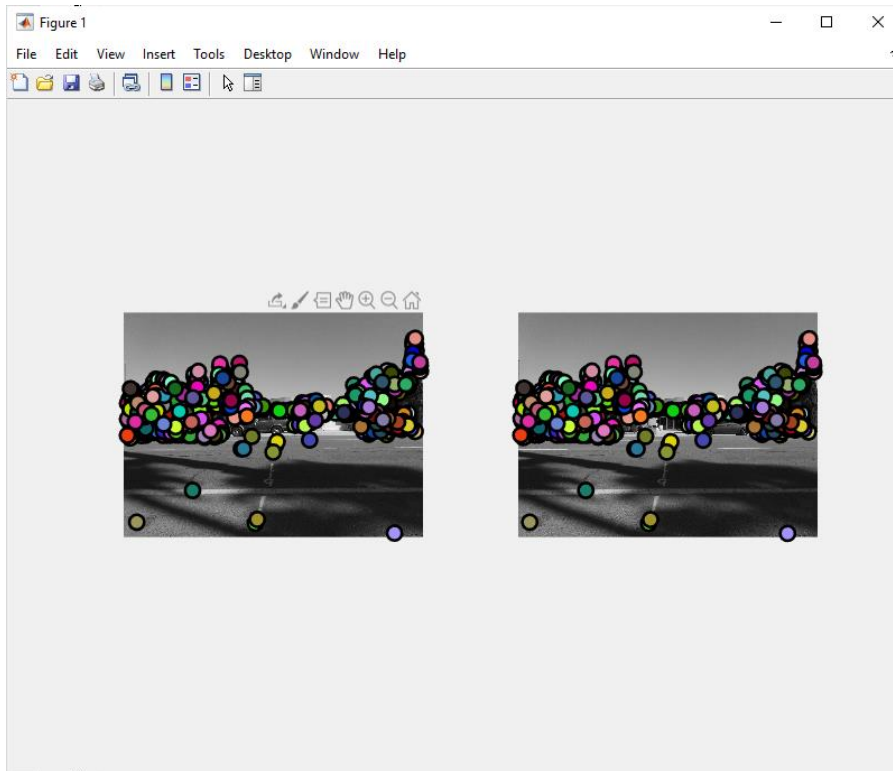
```
23 total good matches, 4 total bad matches
```

```
Saving visualization to eval.jpg
```

The input for some additional examples are as follows:



The output for the same examples are as follows:



Conclusion:

(1) Notre Dame image pair: For case 1, 27 matches were found in total, of which 21 were correct, hinting at an accuracy of about 77% for this one image pair with the matching NNDR ratio threshold at 0.7 and for case 2, 25 matches were found in total, of which 18 were corrected.

(2) CSGO image pair: I have made the ground-truth values and generated the viz.jpg and eval.jpg. For case 1,

23 matches were found in total, of which 16 were corrected, hinting at an accuracy of about 70% for this one image pair with the matching NNDR ratio threshold at 0.7 and for case 2, 27 matches were found in total, of which 23 were corrected.

For image pairs, depending on the set-up of the images, the algorithm can have good or bad performance. As mentioned above, image pairs with large scale or orientation differences are harder to match for the algorithm since it doesn't take these parameters into account. In contrast, image pairs that are similar in scale and orientation are easier to match.